


Neural-Heuristic Pipeline: Propose-and-Repair Architecture

This document formalizes the **Neural-Heuristic Data Pipeline**, a multi-stage workflow that translates raw industrial SKU data into physically stable, high-density warehouse placements. The system employs a "Propose-and-Repair" paradigm, leveraging the strategic intuition of deep learning and the deterministic rigor of physics-based heuristics.

1. The Neural-Heuristic Pipeline Flow

The following flowchart illustrates the end-to-end transition of data from the raw input environment to the finalized, settlement-verified state.



2. Stage 0: Data Decoupling & Payload Ingestion

The pipeline begins by receiving a high-dimensional feature vector representing the state of the SKU and the global environment. To ensure the ML model can make strategic decisions, the system extracts **19 core features**:

Category	Feature Count	Variables
Dimensionality	3	l_i, w_i, h_i (Length, Width, Height)
Physicality	2	m_i (Weight), Fragility_i (Boolean)
Constraints	2	Stackable_i (Boolean), Heavy_i (Boolean)
Environment	6	Warehouse W, D, H , current SKU count, total volume
Orientation	6	Discrete rotation masks (0°, 90°, 180°, etc.)

Transformation: The raw SKU data is normalized and decoupled from the database, feeding the model a standardized strategic context.

3. Stage 1: The Neural Prophet (Strategic Intent)

In this stage, the **Evolutionary Optimizer (EO)** or **Hybrid EO-GA** model acts as the "Prophet." It performs an $O(1)$ inference to propose a spatial coordinate for the i -th item.

- **Objective Function:** $\max f(\hat{\mathbf{x}}_i, \text{State})$
- **Neural Proposal:** The model outputs a raw coordinate set $\hat{\mathbf{p}}_i = [\hat{x}, \hat{y}, \hat{z}, \hat{\theta}]$.

[!NOTE] The Neural stage does not account for collisions or precision. Its role is to provide a "Strategic Hint"—identifying the most mathematically promising zone for space utilization.

4. Stage 4: Deterministic Repair Engine (The Feasibility Layer)

The Heuristic Repair Layer intercepts the neural proposal and projects it onto the nearest feasibility manifold. This resolves the "Neural Drift" where models might propose coordinates that overlap with existing items.


4.1 Intersection-Aware Action Masking

The engine applies a volumetric intersection check against the current occupancy map $\mathcal{M}$: $I(i, j) = \begin{cases} 1 & \text{if } \text{Overlap}(i, j) > 0 \\ 0 & \text{otherwise} \end{cases}$

4.2 Forensic Correction

If $\sum I(i, j) > 0$, the heuristic performs a localized search around the neural anchor $\hat{\mathbf{p}}_i$.

- **Projection:** $\mathbf{x}_{\text{phys}} = \text{clamp}(\hat{\mathbf{x}} \odot \mathbf{W}, 0, \mathbf{W} - \mathbf{d}_i)$.
- **Projection:** $\mathbf{x}_{\text{phys}} = \text{clamp}(\hat{\mathbf{x}} \odot \mathbf{W}, 0, \mathbf{W} - \mathbf{d}_i)$.



Graph 1: The Neural-Heuristic Pipeline Progression — Stage 1: Floating Intent, Stage 2: Feasibility Repair, Stage 3: Physics Settlement.

5. Stage 3: Physics Settlement Layer (The Stability Gate)


The final stage ensures the item respects the laws of gravity. Even a collision-free placement is discarded if it lacks surface support.

5.1 Gravity-Collapse Calculation

The z -coordinate is "collapsed" onto the highest available support surface below the item's projected footprint $[x_i, x_i+l_i] \times [y_i, y_i+w_i]$: $z_{i, \text{final}} = \max(\{0\} \cup \{z_j + h_j \mid \text{Support}(i, j)\})$

5.2 The Stability Gate (SSR)

The Support Stability Rate (SSR) is calculated: $\text{SSR} = \frac{\text{Area}(\text{contact})}{\text{Area}(\text{base})} \geq \text{Threshold}$ (e.g., 0.8)



6. Experimental Validation (Pipeline Efficacy)

The efficacy of this pipeline was validated across three operational scales, comparing standalone heuristics against the neural-hybrid approach.

6.1 Multi-Scale Scaling Analysis


Scale (SKUs)	Algorithm	Repair Latency (ms)	Fitness Score	PSR / SSR	Volume Util (VU)
200	Standalone EO	6,966	30.82%	95.50% / 100%	1.10%
200	Hybrid EO-GA	7,499	31.07%	95.50% / 100%	1.10%
200	Standalone GA	6,696	30.69%	94.00% / 100%	1.07%
200	Hybrid GA-EO	6,805	31.00%	95.50% / 100%	1.09%
400	Standalone EO	34,223	30.69%	94.75% / 100%	2.23%

Scale (SKUs)	Algorithm	Repair Latency (ms)	Fitness Score	PSR / SSR	Volume Util (VU)
400	Hybrid EO-GA	39,061	31.12%	96.50% / 100%	2.26%
400	Standalone GA	37,405	30.77%	97.00% / 100%	2.27%
400	Hybrid GA-EO	34,818	30.97%	95.00% / 100%	2.24%
600	Standalone EO	104,024	30.73%	96.17% / 100%	3.29%
600	Hybrid EO-GA	103,439	31.24%	95.33% / 100%	3.28%
600	Standalone GA	110,575	30.78%	94.83% / 100%	3.24%
600	Hybrid GA-EO	109,287	31.02%	94.50% / 100%	3.25%

6.2 Scaling Metrics & Benchmarks


To visualize the pipeline's behavior as item count increases, we analyze the **Volumetric Utility (VU)**, **Placement Success (PSR)**, and **Inference Throughput**.

Volumetric Scaling



Graph 3: Linear VU Scaling Analysis

PSR Benchmarking





Graph 4: PSR Consistency across Models

Inference Latency

Stability Distribution

Inference Latency

Stability Distribution



Graph 2: Algorithmic Scaling Efficiency

Graph 3: Spatial Support Heatmap (Warehouse Floor)

6.3 Comparison with SOTA (Propose-and-Repair)

Methodology	Support Enforcement	Complexity	Strategic Insight
Classical EP (Crainic et al.)	No	High	None
Action Masking (Zhao et al.)	Partial	Medium	Static
Neural-Heuristic (Ours)	Full (SSR)	Low	Dynamic Evolution

7. Conclusion: The "Strategic Reduction" Benefit

By structuring the system as a pipeline starting with **Neural Intent**, we reduced the heuristic search space from **~4,500 candidates** down to **~120 potential refinements** per item. This hybrid approach enables industrial-scale bin packing with real-time stability guarantees.